

Neural Networks - Project 1

Pegah Givehchian 99222089

Question 1

L1 vs L2 regularization:

In general regularization is used to minimize our loss function and avoid over/under fitting. In both methods we add a value to our loss function to reduce the variance between the train and test datasets and add a little bias.

In [28]: 1 Image("img/L1.png")

Out[28]:
$$\text{LossFunction} = \frac{1}{N} \sum_{i=1}^N (\hat{Y} - Y)^2 + \lambda \sum_{i=1}^N |\theta_i|$$

In [30]: 1 Image("img/L2.png")

Out[30]:
$$\text{LossFunction} = \frac{1}{N} \sum_{i=1}^N (\hat{Y} - Y)^2 + \lambda \sum_{i=1}^N \theta_i^2$$

L1 regularization penalizes the sum of absolute values of the weights, whereas L2 regularization penalizes the sum of squares of the weights.

L1 creates a more sparse model than L2. It has a better performance with outlier points and is more robust, meaning that output and forecast are consistently accurate. It creates a more simple method, which makes it a little harder to learn complicated algorithms. But L2 is less computationally expensive, because L1 takes the absolute values of the weights.

By adding these regularizations, the weights are reduced, so that the models become more simple. So it decreases the chance of overfitting the training data.

Where L1 regularization attempts to estimate the median of data, L2 regularization makes estimation for the mean of the data in order to evade overfitting.

Due to the reasons mentioned above, I would prefer L2 regression. Also while training the models in exercise 2, the model in which I used L2 regression had a better loss than L1 and a lower variance.

Question 2

In [4]:

```

1 import tensorflow as tf
2 import pandas as pd
3 import numpy as np
4 import os
5 import matplotlib.pyplot as plt
6 from tensorflow.keras.models import Sequential
7 from tensorflow.keras.layers import Dense
8 from tensorflow.keras.layers import Dense, BatchNormalization
9 from tensorflow.keras.layers import Dropout
10 from tensorflow.keras.callbacks import EarlyStopping
11 from IPython.display import Image
12 from tensorflow.keras import regularizers

```

Import datasets

In [11]:

```

1 train_data = pd.read_csv("Dataset/fashion-mnist_train.csv")
2 test_data = pd.read_csv("Dataset/fashion-mnist_test.csv")

```

In [12]:

```
1 train_data
```

Out[12]:

	label	pixel1	pixel2	pixel3	pixel4	pixel5	pixel6	pixel7	pixel8	pixel9	...	pixel775	pixe
0	2	0	0	0	0	0	0	0	0	0	0	0	0
1	9	0	0	0	0	0	0	0	0	0	0	0	0
2	6	0	0	0	0	0	0	0	0	5	0	0	0
3	0	0	0	0	1	2	0	0	0	0	0	0	3
4	3	0	0	0	0	0	0	0	0	0	0	0	0
...
59995	9	0	0	0	0	0	0	0	0	0	0	0	0
59996	1	0	0	0	0	0	0	0	0	0	0	0	73
59997	8	0	0	0	0	0	0	0	0	0	0	0	160
59998	8	0	0	0	0	0	0	0	0	0	0	0	0
59999	7	0	0	0	0	0	0	0	0	0	0	0	0

60000 rows × 785 columns



In [13]:

```
1 train_data.shape
```

Out[13]:

```
(60000, 785)
```

```
In [14]: 1 train_data.columns
```

```
Out[14]: Index(['label', 'pixel1', 'pixel2', 'pixel3', 'pixel4', 'pixel5', 'pixel6',
   'pixel7', 'pixel8', 'pixel9',
   ...,
   'pixel775', 'pixel776', 'pixel777', 'pixel778', 'pixel779', 'pixel780',
   'pixel781', 'pixel782', 'pixel783', 'pixel784'],
  dtype='object', length=785)
```

Modeling Simple Model

```
In [15]: 1 trainDataMatrix = np.asmatrix(train_data)
```

```
In [16]: 1 trainDataMatrix
```

```
Out[16]: matrix([[2, 0, 0, ..., 0, 0, 0],
   [9, 0, 0, ..., 0, 0, 0],
   [6, 0, 0, ..., 0, 0, 0],
   ...,
   [8, 0, 0, ..., 0, 0, 0],
   [8, 0, 0, ..., 0, 0, 0],
   [7, 0, 0, ..., 0, 0, 0]], dtype=int64)
```

```
In [17]: 1 Y = trainDataMatrix[:,0]
```

```
In [18]: 1 Y
```

```
Out[18]: matrix([[2],
   [9],
   [6],
   ...,
   [8],
   [8],
   [7]], dtype=int64)
```

```
In [19]: 1 X = trainDataMatrix[:,1:]
```

```
In [20]: 1 X
```

```
Out[20]: matrix([[0, 0, 0, ..., 0, 0, 0],
   [0, 0, 0, ..., 0, 0, 0],
   [0, 0, 0, ..., 0, 0, 0],
   ...,
   [0, 0, 0, ..., 0, 0, 0],
   [0, 0, 0, ..., 0, 0, 0],
   [0, 0, 0, ..., 0, 0, 0]], dtype=int64)
```

```
In [21]: 1 simpleModel = Sequential(  
2     [  
3         tf.keras.Input(shape=(784,)),  
4         Dense(units=225, activation=tf.nn.relu, name="hidden_layer_1"),  
5         Dense(units=256, activation=tf.nn.relu, name="hidden_layer_2"),  
6         Dense(units=128, activation=tf.nn.relu, name="hidden_layer_3"),  
7         Dense(units=10, activation=tf.nn.softmax, name="output_layer"),  
8     ], name = "TheFirstModel"  
9 )
```

```
In [15]: 1 simpleModel.compile(  
2     loss = "sparse_categorical_crossentropy",  
3     optimizer = tf.keras.optimizers.Adam(0.001),  
4 )  
5 simpleModel.fit(  
6     X, Y,  
7     epochs=15  
8 )
```

```
Epoch 1/15  
1875/1875 [=====] - 9s 4ms/step - loss: 1.3489  
Epoch 2/15  
1875/1875 [=====] - 8s 4ms/step - loss: 0.4785  
Epoch 3/15  
1875/1875 [=====] - 9s 5ms/step - loss: 0.4363  
Epoch 4/15  
1875/1875 [=====] - 8s 5ms/step - loss: 0.4250  
Epoch 5/15  
1875/1875 [=====] - 8s 4ms/step - loss: 0.3976  
Epoch 6/15  
1875/1875 [=====] - 14s 8ms/step - loss: 0.3775  
Epoch 7/15  
1875/1875 [=====] - 12s 7ms/step - loss: 0.3615  
Epoch 8/15  
1875/1875 [=====] - 9s 5ms/step - loss: 0.3529  
Epoch 9/15  
1875/1875 [=====] - 8s 5ms/step - loss: 0.3449  
Epoch 10/15  
1875/1875 [=====] - 9s 5ms/step - loss: 0.3332  
Epoch 11/15  
1875/1875 [=====] - 9s 5ms/step - loss: 0.3340  
Epoch 12/15  
1875/1875 [=====] - 8s 4ms/step - loss: 0.3195  
Epoch 13/15  
1875/1875 [=====] - 8s 4ms/step - loss: 0.3147  
Epoch 14/15  
1875/1875 [=====] - 9s 5ms/step - loss: 0.3093  
Epoch 15/15  
1875/1875 [=====] - 8s 4ms/step - loss: 0.3027
```

```
Out[15]: <keras.callbacks.History at 0x2b6139919a0>
```

```
In [24]: 1 testMatrix = np.asmatrix(test_data)
          2 xTest = testMatrix[:,1:]
          3 yTest = testMatrix[:,0]
```

```
In [17]: 1 simpleModel.evaluate(xTest, yTest)
```

```
313/313 [=====] - 1s 2ms/step - loss: 0.3974
```

```
Out[17]: 0.3973921239376068
```

We try to decrease the variance and loss by implementing different methods and regularizations.

Depth Effect

The number of layers in a neural network defines its depth. So I try to create different models with different number of layers to understand its effect.

5 Layers

```
In [72]: 1 count5LayerModel = Sequential(
          2 [
          3     tf.keras.Input(shape=(784,)),
          4     Dense(units=225, activation=tf.nn.relu, name="hidden_layer_1"),
          5     Dense(units=256, activation=tf.nn.relu, name="hidden_layer_2"),
          6     Dense(units=128, activation=tf.nn.relu, name="hidden_layer_3"),
          7     Dense(units=128, activation=tf.nn.relu, name="hidden_layer_4"),
          8     Dense(units=128, activation=tf.nn.relu, name="hidden_layer_5"),
          9     Dense(units=10, activation=tf.nn.softmax, name="output_layer"),
         10 ], name = "count5LayerModel"
         11 )
```

In [73]:

```
1 count5LayerModel.compile(  
2     loss = "sparse_categorical_crossentropy",  
3     optimizer = tf.keras.optimizers.Adam(0.001),  
4 )  
5 count5LayerModel.fit(  
6     X,Y,  
7     epochs=15  
8 )
```

```
Epoch 1/15  
1875/1875 [=====] - 8s 4ms/step - loss: 0.8956  
Epoch 2/15  
1875/1875 [=====] - 8s 4ms/step - loss: 0.4703  
Epoch 3/15  
1875/1875 [=====] - 7s 4ms/step - loss: 0.4269  
Epoch 4/15  
1875/1875 [=====] - 7s 4ms/step - loss: 0.4012  
Epoch 5/15  
1875/1875 [=====] - 7s 4ms/step - loss: 0.3773  
Epoch 6/15  
1875/1875 [=====] - 6s 3ms/step - loss: 0.3627  
Epoch 7/15  
1875/1875 [=====] - 7s 4ms/step - loss: 0.3471  
Epoch 8/15  
1875/1875 [=====] - 7s 4ms/step - loss: 0.3332  
Epoch 9/15  
1875/1875 [=====] - 7s 4ms/step - loss: 0.3258  
Epoch 10/15  
1875/1875 [=====] - 7s 3ms/step - loss: 0.3216  
Epoch 11/15  
1875/1875 [=====] - 7s 4ms/step - loss: 0.3121  
Epoch 12/15  
1875/1875 [=====] - 7s 4ms/step - loss: 0.2979  
Epoch 13/15  
1875/1875 [=====] - 7s 4ms/step - loss: 0.2955  
Epoch 14/15  
1875/1875 [=====] - 8s 4ms/step - loss: 0.2952  
Epoch 15/15  
1875/1875 [=====] - 7s 4ms/step - loss: 0.2914
```

Out[73]: <keras.callbacks.History at 0x2b6142b7a60>

In [74]:

```
1 count5LayerModel.evaluate(xTest, yTest)
```

```
313/313 [=====] - 1s 3ms/step - loss: 0.4297
```

Out[74]: 0.4297083914279938

2 Layers

```
In [75]: 1 count2LayerModel = Sequential(  
2     [  
3         tf.keras.Input(shape=(784,)),  
4         Dense(units=225, activation=tf.nn.relu, name="hidden_layer_1"),  
5         Dense(units=256, activation=tf.nn.relu, name="hidden_layer_2"),  
6         Dense(units=10, activation=tf.nn.softmax, name="output_layer"),  
7     ], name = "count2LayerModel"  
8 )
```

```
In [76]: 1 count2LayerModel.compile(  
2     loss = "sparse_categorical_crossentropy",  
3     optimizer = tf.keras.optimizers.Adam(0.001),  
4 )  
5 count2LayerModel.fit(  
6     X,Y,  
7     epochs=15  
8 )
```

```
Epoch 1/15  
1875/1875 [=====] - 6s 3ms/step - loss: 2.1136  
Epoch 2/15  
1875/1875 [=====] - 5s 3ms/step - loss: 0.6514  
Epoch 3/15  
1875/1875 [=====] - 5s 3ms/step - loss: 0.5859  
Epoch 4/15  
1875/1875 [=====] - 5s 3ms/step - loss: 0.5229  
Epoch 5/15  
1875/1875 [=====] - 6s 3ms/step - loss: 0.4697  
Epoch 6/15  
1875/1875 [=====] - 5s 3ms/step - loss: 0.4430  
Epoch 7/15  
1875/1875 [=====] - 5s 3ms/step - loss: 0.4218  
Epoch 8/15  
1875/1875 [=====] - 5s 3ms/step - loss: 0.4088  
Epoch 9/15  
1875/1875 [=====] - 5s 3ms/step - loss: 0.4059  
Epoch 10/15  
1875/1875 [=====] - 5s 3ms/step - loss: 0.3915  
Epoch 11/15  
1875/1875 [=====] - 5s 3ms/step - loss: 0.3835  
Epoch 12/15  
1875/1875 [=====] - 6s 3ms/step - loss: 0.3791  
Epoch 13/15  
1875/1875 [=====] - 5s 3ms/step - loss: 0.3793  
Epoch 14/15  
1875/1875 [=====] - 5s 3ms/step - loss: 0.3770  
Epoch 15/15  
1875/1875 [=====] - 6s 3ms/step - loss: 0.3681
```

```
Out[76]: <keras.callbacks.History at 0x2b6143da7c0>
```

In [77]: 1 count2LayerModel.evaluate(xTest, yTest)

```
313/313 [=====] - 1s 2ms/step - loss: 0.4555
```

Out[77]: 0.45554590225219727

10 Layers

In [78]: 1 count10LayerModel = Sequential(
2 [
3 tf.keras.Input(shape=(784,)),
4 Dense(units=225, activation=tf.nn.relu, name="hidden_layer_1"),
5 Dense(units=256, activation=tf.nn.relu, name="hidden_layer_2"),
6 Dense(units=128, activation=tf.nn.relu, name="hidden_layer_3"),
7 Dense(units=128, activation=tf.nn.relu, name="hidden_layer_4"),
8 Dense(units=225, activation=tf.nn.relu, name="hidden_layer_5"),
9 Dense(units=256, activation=tf.nn.relu, name="hidden_layer_6"),
10 Dense(units=28, activation=tf.nn.relu, name="hidden_layer_7"),
11 Dense(units=64, activation=tf.nn.relu, name="hidden_layer_8"),
12 Dense(units=128, activation=tf.nn.relu, name="hidden_layer_9"),
13 Dense(units=225, activation=tf.nn.relu, name="hidden_layer_10"),
14 Dense(units=10, activation=tf.nn.softmax, name="output_layer"),
15], name = "count10LayerModel"
16)

In [79]:

```

1 count10LayerModel.compile(
2     loss = "sparse_categorical_crossentropy",
3     optimizer = tf.keras.optimizers.Adam(0.001),
4 )
5 count10LayerModel.fit(
6     X,Y,
7     epochs=15
8 )

```

```

Epoch 1/15
1875/1875 [=====] - 9s 4ms/step - loss: 0.6390
Epoch 2/15
1875/1875 [=====] - 10s 5ms/step - loss: 0.4669
Epoch 3/15
1875/1875 [=====] - 10s 5ms/step - loss: 0.4246
Epoch 4/15
1875/1875 [=====] - 9s 5ms/step - loss: 0.3982
Epoch 5/15
1875/1875 [=====] - 9s 5ms/step - loss: 0.3757
Epoch 6/15
1875/1875 [=====] - 9s 5ms/step - loss: 0.3686
Epoch 7/15
1875/1875 [=====] - 9s 5ms/step - loss: 0.3584
Epoch 8/15
1875/1875 [=====] - 9s 5ms/step - loss: 0.3445
Epoch 9/15
1875/1875 [=====] - 9s 5ms/step - loss: 0.3393
Epoch 10/15
1875/1875 [=====] - 9s 5ms/step - loss: 0.3391
Epoch 11/15
1875/1875 [=====] - 9s 5ms/step - loss: 0.3245
Epoch 12/15
1875/1875 [=====] - 9s 5ms/step - loss: 0.3205
Epoch 13/15
1875/1875 [=====] - 9s 5ms/step - loss: 0.3200
Epoch 14/15
1875/1875 [=====] - 9s 5ms/step - loss: 0.3074
Epoch 15/15
1875/1875 [=====] - 9s 5ms/step - loss: 0.3063

```

Out[79]: <keras.callbacks.History at 0x2b615bee850>

In [80]:

```
1 count10LayerModel.evaluate(xTest, yTest)
```

```
313/313 [=====] - 1s 3ms/step - loss: 0.4052
```

Out[80]: 0.40516334772109985

We see that the more depth we add to our network, the less our loss becomes, but our variance becomes higher, which means that our model is overfitting.

Normalization

We devide our inputs by 255 so that they are between 0 and 1

In [18]: 1 normalX = X/255

In [19]: 1 simpleModel.compile(
2 loss = "sparse_categorical_crossentropy",
3 optimizer = tf.keras.optimizers.Adam(0.001),
4)
5 simpleModel.fit(
6 normalX,Y,
7 epochs=15
8)

```
Epoch 1/15
1875/1875 [=====] - 8s 4ms/step - loss: 0.3547
Epoch 2/15
1875/1875 [=====] - 8s 4ms/step - loss: 0.2797
Epoch 3/15
1875/1875 [=====] - 8s 4ms/step - loss: 0.2634
Epoch 4/15
1875/1875 [=====] - 8s 4ms/step - loss: 0.2512
Epoch 5/15
1875/1875 [=====] - 8s 4ms/step - loss: 0.2429
Epoch 6/15
1875/1875 [=====] - 8s 4ms/step - loss: 0.2352
Epoch 7/15
1875/1875 [=====] - 8s 4ms/step - loss: 0.2269
Epoch 8/15
1875/1875 [=====] - 10s 5ms/step - loss: 0.2206
Epoch 9/15
1875/1875 [=====] - 9s 5ms/step - loss: 0.2147
Epoch 10/15
1875/1875 [=====] - 8s 4ms/step - loss: 0.2067
Epoch 11/15
1875/1875 [=====] - 8s 4ms/step - loss: 0.2009
Epoch 12/15
1875/1875 [=====] - 12s 6ms/step - loss: 0.1971
Epoch 13/15
1875/1875 [=====] - 10s 5ms/step - loss: 0.1905
Epoch 14/15
1875/1875 [=====] - 8s 4ms/step - loss: 0.1866
Epoch 15/15
1875/1875 [=====] - 8s 4ms/step - loss: 0.1806
```

Out[19]: <keras.callbacks.History at 0x2b6137c1190>

In [20]: 1 simpleModel.evaluate(xTest/255, yTest)

```
313/313 [=====] - 1s 2ms/step - loss: 0.3517
```

Out[20]: 0.3516875207424164

we see that our loss has decreased, but we have a high variance, so our model is still overfitting.

Batch Normalization

```
In [21]: 1 Model2 = Sequential(  
2     [  
3         tf.keras.Input(shape=(784,)),  
4         Dense(units=225, activation=tf.nn.relu, name="hidden_layer_1"),  
5         BatchNormalization(),  
6         Dense(units=256, activation=tf.nn.relu, name="hidden_layer_2"),  
7         BatchNormalization(),  
8         Dense(units=128, activation=tf.nn.relu, name="hidden_layer_3"),  
9         BatchNormalization(),  
10        Dense(units=10, activation=tf.nn.softmax, name="output_layer"),  
11    ], name = "BatchNormalizationModel"  
12 )
```

In [23]:

```
1 Model2.compile(  
2     loss = "sparse_categorical_crossentropy",  
3     optimizer = tf.keras.optimizers.Adam(0.001),  
4 )  
5 Model2.fit(  
6     X,Y,  
7     epochs=15  
8 )
```

```
Epoch 1/15  
1875/1875 [=====] - 11s 5ms/step - loss: 0.4994  
Epoch 2/15  
1875/1875 [=====] - 10s 5ms/step - loss: 0.3936  
Epoch 3/15  
1875/1875 [=====] - 10s 5ms/step - loss: 0.3571  
Epoch 4/15  
1875/1875 [=====] - 10s 5ms/step - loss: 0.3346  
Epoch 5/15  
1875/1875 [=====] - 10s 5ms/step - loss: 0.3200  
Epoch 6/15  
1875/1875 [=====] - 10s 5ms/step - loss: 0.3026  
Epoch 7/15  
1875/1875 [=====] - 10s 6ms/step - loss: 0.2917  
Epoch 8/15  
1875/1875 [=====] - 10s 5ms/step - loss: 0.2795  
Epoch 9/15  
1875/1875 [=====] - 10s 5ms/step - loss: 0.2715  
Epoch 10/15  
1875/1875 [=====] - 11s 6ms/step - loss: 0.2642  
Epoch 11/15  
1875/1875 [=====] - 12s 7ms/step - loss: 0.2531  
Epoch 12/15  
1875/1875 [=====] - 12s 6ms/step - loss: 0.2432  
Epoch 13/15  
1875/1875 [=====] - 11s 6ms/step - loss: 0.2366  
Epoch 14/15  
1875/1875 [=====] - 17s 9ms/step - loss: 0.2302  
Epoch 15/15  
1875/1875 [=====] - 14s 7ms/step - loss: 0.2221
```

Out[23]: <keras.callbacks.History at 0x2b616611c10>

In [31]:

```
1 Model2.evaluate(xTest, yTest)
```

```
313/313 [=====] - 1s 2ms/step - loss: 0.4204
```

Out[31]: 0.420440137386322

Batch with the use of normalization

In [32]:

```

1 Model2.compile(
2     loss = "sparse_categorical_crossentropy",
3     optimizer = tf.keras.optimizers.Adam(0.001),
4 )
5 Model2.fit(
6     normalX, Y,
7     epochs=15
8 )

```

```

Epoch 1/15
1875/1875 [=====] - 10s 4ms/step - loss: 0.2197
Epoch 2/15
1875/1875 [=====] - 9s 5ms/step - loss: 0.2140
Epoch 3/15
1875/1875 [=====] - 9s 5ms/step - loss: 0.2069
Epoch 4/15
1875/1875 [=====] - 9s 5ms/step - loss: 0.2026
Epoch 5/15
1875/1875 [=====] - 9s 5ms/step - loss: 0.1986
Epoch 6/15
1875/1875 [=====] - 10s 5ms/step - loss: 0.1948
Epoch 7/15
1875/1875 [=====] - 10s 5ms/step - loss: 0.1893
Epoch 8/15
1875/1875 [=====] - 10s 5ms/step - loss: 0.1838
Epoch 9/15
1875/1875 [=====] - 9s 5ms/step - loss: 0.1810
Epoch 10/15
1875/1875 [=====] - 9s 5ms/step - loss: 0.1766
Epoch 11/15
1875/1875 [=====] - 9s 5ms/step - loss: 0.1730
Epoch 12/15
1875/1875 [=====] - 9s 5ms/step - loss: 0.1686
Epoch 13/15
1875/1875 [=====] - 9s 5ms/step - loss: 0.1673
Epoch 14/15
1875/1875 [=====] - 8s 4ms/step - loss: 0.1687
Epoch 15/15
1875/1875 [=====] - 8s 5ms/step - loss: 0.1631

```

Out[32]: <keras.callbacks.History at 0x2b616690be0>

In [33]:

```
1 Model2.evaluate(xTest/255, yTest)
```

```
313/313 [=====] - 1s 2ms/step - loss: 0.3373
```

Out[33]: 0.33728882670402527

we see that batch normalization's loss is less than the simple model. But it has a high variance because the loss in the test and train datasets differ a lot, which means our model is yet still overfitting, so we use different methods like the dropout technique, early stopping or L1 L2 regularizations to decrease the variance and avoid overfitting.

Dropout

Dropout is generally used for weakening the model to avoid overfitting, by dropping randomly selected neurons.

In [60]:

```
1 Model3 = Sequential(
2     [
3         tf.keras.Input(shape=(784,)),
4         Dense(units=225, activation=tf.nn.relu, name="hidden_layer_1"),
5         Dropout(0.3),
6         Dense(units=256, activation=tf.nn.relu, name="hidden_layer_2"),
7         Dropout(0.5),
8         Dense(units=128, activation=tf.nn.relu, name="hidden_layer_3"),
9         Dense(units=10, activation=tf.nn.softmax, name="output_layer"),
10    ], name = "DropOutModel"
11 )
```

In [61]:

```

1 Model3.compile(
2     loss = "sparse_categorical_crossentropy",
3     optimizer = tf.keras.optimizers.Adam(0.001),
4 )
5 Model3.fit(
6     X,Y,
7     epochs=15
8 )

```

```

Epoch 1/15
1875/1875 [=====] - 6s 3ms/step - loss: 2.4973
Epoch 2/15
1875/1875 [=====] - 6s 3ms/step - loss: 1.0378
Epoch 3/15
1875/1875 [=====] - 6s 3ms/step - loss: 0.8729
Epoch 4/15
1875/1875 [=====] - 7s 4ms/step - loss: 0.8187
Epoch 5/15
1875/1875 [=====] - 7s 4ms/step - loss: 0.7598
Epoch 6/15
1875/1875 [=====] - 7s 4ms/step - loss: 0.7460
Epoch 7/15
1875/1875 [=====] - 7s 4ms/step - loss: 0.7254
Epoch 8/15
1875/1875 [=====] - 7s 4ms/step - loss: 0.7197
Epoch 9/15
1875/1875 [=====] - 6s 3ms/step - loss: 0.6960
Epoch 10/15
1875/1875 [=====] - 7s 4ms/step - loss: 0.6907
Epoch 11/15
1875/1875 [=====] - 7s 4ms/step - loss: 0.6854
Epoch 12/15
1875/1875 [=====] - 6s 3ms/step - loss: 0.6825
Epoch 13/15
1875/1875 [=====] - 7s 4ms/step - loss: 0.6894
Epoch 14/15
1875/1875 [=====] - 8s 4ms/step - loss: 0.6654
Epoch 15/15
1875/1875 [=====] - 8s 4ms/step - loss: 0.6763

```

Out[61]: <keras.callbacks.History at 0x2b61523e490>

In [62]:

```

1 Model3.evaluate(xTest, yTest)

```

```

313/313 [=====] - 1s 2ms/step - loss: 0.5810

```

Out[62]: 0.5810006856918335

We see that adding dropout gives us a higher loss than the simple model, but a lower variance, which makes our model more trustworthy.

Early Stopping Criteria

One of the methods used to avoid overfitting. It stops training as soon as the validation error reaches a minimum.

In [63]:

```
1 Model4 = Sequential(
2     [
3         tf.keras.Input(shape=(784,)),
4         Dense(units=225 ,activation=tf.nn.relu ,name="hidden_layer_1"),
5         Dense(units=256 ,activation=tf.nn.relu, name="hidden_layer_2"),
6         Dense(units=128 ,activation=tf.nn.relu, name="hidden_layer_3"),
7         Dense(units=10 ,activation=tf.nn.softmax, name="output_layer"),
8     ],name = "EarlyStoppingModel"
9 )
```

In [64]:

```
1 Model4.compile(  
2     loss = "sparse_categorical_crossentropy",  
3     optimizer = tf.keras.optimizers.Adam(0.001),  
4 )  
5 Model4.fit(  
6     X,Y,  
7     epochs=15,  
8     callbacks=[EarlyStopping()])  
9 )
```

```
Epoch 1/15  
1868/1875 [=====>.] - ETA: 0s - loss: 1.3881WARNING:tens  
orflow:Early stopping conditioned on metric `val_loss` which is not available.  
Available metrics are: loss  
1875/1875 [=====] - 5s 2ms/step - loss: 1.3856  
Epoch 2/15  
1869/1875 [=====>.] - ETA: 0s - loss: 0.5191WARNING:tens  
orflow:Early stopping conditioned on metric `val_loss` which is not available.  
Available metrics are: loss  
1875/1875 [=====] - 6s 3ms/step - loss: 0.5186  
Epoch 3/15  
1866/1875 [=====>.] - ETA: 0s - loss: 0.4545WARNING:tens  
orflow:Early stopping conditioned on metric `val_loss` which is not available.  
Available metrics are: loss  
1875/1875 [=====] - 6s 3ms/step - loss: 0.4545  
Epoch 4/15  
1861/1875 [=====>.] - ETA: 0s - loss: 0.4161WARNING:tens  
orflow:Early stopping conditioned on metric `val_loss` which is not available.  
Available metrics are: loss  
1875/1875 [=====] - 6s 3ms/step - loss: 0.4164  
Epoch 5/15  
1868/1875 [=====>.] - ETA: 0s - loss: 0.3996WARNING:tens  
orflow:Early stopping conditioned on metric `val_loss` which is not available.  
Available metrics are: loss  
1875/1875 [=====] - 6s 3ms/step - loss: 0.3996  
Epoch 6/15  
1875/1875 [=====] - ETA: 0s - loss: 0.3790WARNING:tens  
orflow:Early stopping conditioned on metric `val_loss` which is not available.  
Available metrics are: loss  
1875/1875 [=====] - 6s 3ms/step - loss: 0.3790  
Epoch 7/15  
1862/1875 [=====>.] - ETA: 0s - loss: 0.3628WARNING:tens  
orflow:Early stopping conditioned on metric `val_loss` which is not available.  
Available metrics are: loss  
1875/1875 [=====] - 6s 3ms/step - loss: 0.3633  
Epoch 8/15  
1861/1875 [=====>.] - ETA: 0s - loss: 0.3524WARNING:tens  
orflow:Early stopping conditioned on metric `val_loss` which is not available.  
Available metrics are: loss  
1875/1875 [=====] - 7s 4ms/step - loss: 0.3524  
Epoch 9/15  
1870/1875 [=====>.] - ETA: 0s - loss: 0.3457WARNING:tens  
orflow:Early stopping conditioned on metric `val_loss` which is not available.  
Available metrics are: loss  
1875/1875 [=====] - 7s 3ms/step - loss: 0.3457  
Epoch 10/15
```

```
1875/1875 [=====] - ETA: 0s - loss: 0.3358WARNING:tens
orflow:Early stopping conditioned on metric `val_loss` which is not available.
Available metrics are: loss
1875/1875 [=====] - 7s 4ms/step - loss: 0.3358
Epoch 11/15
1866/1875 [=====>.] - ETA: 0s - loss: 0.3365WARNING:tens
orflow:Early stopping conditioned on metric `val_loss` which is not available.
Available metrics are: loss
1875/1875 [=====] - 7s 4ms/step - loss: 0.3373
Epoch 12/15
1870/1875 [=====>.] - ETA: 0s - loss: 0.3258WARNING:tens
orflow:Early stopping conditioned on metric `val_loss` which is not available.
Available metrics are: loss
1875/1875 [=====] - 7s 4ms/step - loss: 0.3258
Epoch 13/15
1872/1875 [=====>.] - ETA: 0s - loss: 0.3236WARNING:tens
orflow:Early stopping conditioned on metric `val_loss` which is not available.
Available metrics are: loss
1875/1875 [=====] - 7s 4ms/step - loss: 0.3240
Epoch 14/15
1875/1875 [=====] - ETA: 0s - loss: 0.3152WARNING:tens
orflow:Early stopping conditioned on metric `val_loss` which is not available.
Available metrics are: loss
1875/1875 [=====] - 7s 4ms/step - loss: 0.3152
Epoch 15/15
1874/1875 [=====>.] - ETA: 0s - loss: 0.3122WARNING:tens
orflow:Early stopping conditioned on metric `val_loss` which is not available.
Available metrics are: loss
1875/1875 [=====] - 7s 4ms/step - loss: 0.3122
```

Out[64]: <keras.callbacks.History at 0x2b6140c3b20>

In [65]: 1 Model4.evaluate(xTest, yTest)

```
313/313 [=====] - 1s 2ms/step - loss: 0.3754
```

Out[65]: 0.3753807246685028

We see that it has a high loss, but it is lower than the dropout technique.

L1 Regularization

In [57]: 1 Model5 = Sequential(
2 [
3 tf.keras.Input(shape=(784,)),
4 Dense(units=225 ,activation=tf.nn.relu ,name="hidden_layer_1", kerne,
5 Dense(units=256 ,activation=tf.nn.relu,name="hidden_layer_2", kernel,
6 Dense(units=128 ,activation=tf.nn.relu,name="hidden_layer_3", kernel,
7 Dense(units=10 ,activation=tf.nn.softmax,name="output_layer"),
8],name = "L1Model"
9)

In [58]:

```

1 Model5.compile(
2     loss = "sparse_categorical_crossentropy",
3     optimizer = tf.keras.optimizers.Adam(0.001),
4 )
5 Model5.fit(
6     X,Y,
7     epochs=15,
8 )

```

```

Epoch 1/15
1875/1875 [=====] - 8s 4ms/step - loss: 33.6343
Epoch 2/15
1875/1875 [=====] - 8s 4ms/step - loss: 2.4071
Epoch 3/15
1875/1875 [=====] - 7s 4ms/step - loss: 1.2655
Epoch 4/15
1875/1875 [=====] - 7s 4ms/step - loss: 1.2142
Epoch 5/15
1875/1875 [=====] - 7s 4ms/step - loss: 1.2271
Epoch 6/15
1875/1875 [=====] - 7s 4ms/step - loss: 1.2136
Epoch 7/15
1875/1875 [=====] - 8s 4ms/step - loss: 1.2111
Epoch 8/15
1875/1875 [=====] - 7s 4ms/step - loss: 1.2010
Epoch 9/15
1875/1875 [=====] - 8s 4ms/step - loss: 1.1891
Epoch 10/15
1875/1875 [=====] - 8s 4ms/step - loss: 1.2136
Epoch 11/15
1875/1875 [=====] - 9s 5ms/step - loss: 1.1858
Epoch 12/15
1875/1875 [=====] - 9s 5ms/step - loss: 1.1924
Epoch 13/15
1875/1875 [=====] - 8s 4ms/step - loss: 1.2010
Epoch 14/15
1875/1875 [=====] - 9s 5ms/step - loss: 1.2041
Epoch 15/15
1875/1875 [=====] - 8s 4ms/step - loss: 1.1852

```

Out[58]: <keras.callbacks.History at 0x2b61470a190>

In [66]:

```
1 Model5.evaluate(xTest, yTest)
```

```
313/313 [=====] - 1s 2ms/step - loss: 1.1481
```

Out[66]: 1.1480625867843628

This regularization technique gives us an even higher loss than the dropout or early stopping methods, but it's lesss overfitting.

L2 Regularization

```
In [67]: 1 Model6 = Sequential(  
2     [  
3         tf.keras.Input(shape=(784,)),  
4         Dense(units=225, activation=tf.nn.relu, name="hidden_layer_1", kernel  
5             Dense(units=256, activation=tf.nn.relu, name="hidden_layer_2", kernel  
6                 Dense(units=128, activation=tf.nn.relu, name="hidden_layer_3", kernel  
7                     Dense(units=10, activation=tf.nn.softmax, name="output_layer"),  
8                 ], name="L2Model"  
9     )
```

```
In [68]: 1 Model6.compile(  
2     loss="sparse_categorical_crossentropy",  
3     optimizer=tf.keras.optimizers.Adam(0.001),  
4 )  
5 Model6.fit(  
6     X, Y,  
7     epochs=15,  
8 )
```

```
Epoch 1/15  
1875/1875 [=====] - 7s 4ms/step - loss: 5.8678  
Epoch 2/15  
1875/1875 [=====] - 6s 3ms/step - loss: 2.0531  
Epoch 3/15  
1875/1875 [=====] - 6s 3ms/step - loss: 0.9369  
Epoch 4/15  
1875/1875 [=====] - 6s 3ms/step - loss: 0.6412  
Epoch 5/15  
1875/1875 [=====] - 6s 3ms/step - loss: 0.5741  
Epoch 6/15  
1875/1875 [=====] - 6s 3ms/step - loss: 0.5562  
Epoch 7/15  
1875/1875 [=====] - 6s 3ms/step - loss: 0.5541  
Epoch 8/15  
1875/1875 [=====] - 6s 3ms/step - loss: 0.5450  
Epoch 9/15  
1875/1875 [=====] - 6s 3ms/step - loss: 0.5356  
Epoch 10/15  
1875/1875 [=====] - 7s 4ms/step - loss: 0.5355  
Epoch 11/15  
1875/1875 [=====] - 7s 4ms/step - loss: 0.5293  
Epoch 12/15  
1875/1875 [=====] - 6s 3ms/step - loss: 0.5277  
Epoch 13/15  
1875/1875 [=====] - 6s 3ms/step - loss: 0.5286  
Epoch 14/15  
1875/1875 [=====] - 7s 4ms/step - loss: 0.5236  
Epoch 15/15  
1875/1875 [=====] - 7s 4ms/step - loss: 0.5213
```

```
Out[68]: <keras.callbacks.History at 0x2b6150adb0>
```

In [69]: 1 Model6.evaluate(xTest, yTest)

```
313/313 [=====] - 1s 2ms/step - loss: 0.5247
```

Out[69]: 0.5246573686599731

L2's loss is less than L1. Variance is very low, because the loss in the train and test dataset doesn't differ a lot, Which means L2 is a good regularizer.

L1_L2 Regularization

In [5]: 1 Model7 = Sequential(
2 [
3 tf.keras.Input(shape=(784,)),
4 Dense(units=225, activation=tf.nn.relu, name="hidden_layer_1", kernel
5 Dense(units=256, activation=tf.nn.relu, name="hidden_layer_2", kernel
6 Dense(units=128, activation=tf.nn.relu, name="hidden_layer_3", kernel
7 Dense(units=10, activation=tf.nn.softmax, name="output_layer"),
8], name = "L1_L2Model"
9)

In [22]: 1 Model7.compile(
2 loss = "sparse_categorical_crossentropy",
3 optimizer = tf.keras.optimizers.Adam(0.001),
4)
5 Model7.fit(
6 X, Y,
7 epochs=15,
8)

```
Epoch 1/15
1875/1875 [=====] - 12s 6ms/step - loss: 29.9522
Epoch 2/15
1875/1875 [=====] - 11s 6ms/step - loss: 2.1711
Epoch 3/15
1875/1875 [=====] - 11s 6ms/step - loss: 1.3294
Epoch 4/15
1875/1875 [=====] - 11s 6ms/step - loss: 1.2652
Epoch 5/15
1875/1875 [=====] - 11s 6ms/step - loss: 1.2577
Epoch 6/15
1875/1875 [=====] - 11s 6ms/step - loss: 1.2485
Epoch 7/15
1875/1875 [=====] - 11s 6ms/step - loss: 1.2376
Epoch 8/15
1875/1875 [=====] - 11s 6ms/step - loss: 1.2239
Epoch 9/15
1875/1875 [=====] - 11s 6ms/step - loss: 1.2035
Epoch 10/15
1875/1875 [=====] - 11s 6ms/step - loss: 1.2121
```

In [25]: 1 Model7.evaluate(xTest, yTest)

```
313/313 [=====] - 2s 4ms/step - loss: 1.2572
```

Out[25]: 1.2572354078292847

we see that the variance is low, but the loss is not good.

In general i think that L2 had the best influence on the model for avoiding overfitting.

In []: 1