

Intel Classification using the resnet architecture

Pegah Givehchian - 99222089 - Seri 3 - Computational neural networks

```
In [1]: 1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 from collections import OrderedDict
6 import cv2
7 import os
8 from PIL import Image
9 import torch
10 from torch import optim
11 from torch.autograd import Variable
12 from torch.utils.data import random_split, DataLoader
13 from torch import nn
14 import torch.nn.functional as F
15 from torchvision.utils import make_grid
16 from torchvision import transforms, models, datasets
17 from tqdm import tqdm, trange, tqdm
```

Define the paths

```
In [2]: 1 train_dir = "../input/intel-image-classification/seg_train/seg_train/"
2 test_dir = "../input/intel-image-classification/seg_test/seg_test/"
3 pred_dir = "../input/intel-image-classification/seg_pred/seg_pred/"
4
5
6 pred_files = [os.path.join(pred_dir, f) for f in os.listdir(pred_dir)]
```

Define our classes

```
In [3]: 1 class_arr = os.listdir(train_dir)
2 print(enumerate([0,1]))
3 classes = {0: 'buildings',
4           1: 'forest',
5           2: 'glacier',
6           3: 'mountain',
7           4: 'sea',
8           5: 'street'}
9 classes
```

<enumerate object at 0x7f99e52d2dc0>

```
Out[3]: {0: 'buildings',
1: 'forest',
2: 'glacier',
3: 'mountain',
4: 'sea',
5: 'street'}
```

Get the count of each category in our training and testing data:

```
In [4]: 1 def get_cat_count(dir):
2     cat_count = {}
3     for cat in class_arr:
4         count = len(os.listdir(os.path.join(dir, cat)))
5         cat_count[cat] = count
6     return(cat_count)
```

```
In [5]: 1 get_cat_count(train_dir)
```

```
Out[5]: {'mountain': 2512,
'street': 2382,
'buildings': 2191,
'sea': 2274,
'forest': 2271,
'glacier': 2404}
```


```
In [6]: get_cat_count(test_dir)
```

```
Out[6]: {'mountain': 525,  
         'street': 501,  
         'buildings': 437,  
         'sea': 510,  
         'forest': 474,  
         'glacier': 553}
```

Transform data

```
In [7]: mean = [0.485, 0.456, 0.406]  
std = [0.229, 0.224, 0.225]  
  
train_transform = transforms.Compose([transforms.Resize((150, 150)),  
                                     transforms.RandomResizedCrop(150),  
                                     transforms.RandomRotation(30),  
                                     transforms.RandomHorizontalFlip(),  
                                     transforms.ToTensor(),  
                                     transforms.Normalize(torch.Tensor(mean), torch.Tensor(std)) # Normalize  
                                     ])  
  
test_transform = transforms.Compose([transforms.Resize((150, 150)),  
                                    transforms.CenterCrop(150),  
                                    transforms.ToTensor(),  
                                    transforms.Normalize(torch.Tensor(mean), torch.Tensor(std))  
                                    ])  
  
tmp_ds = datasets.ImageFolder(train_dir, transform=train_transform)  
  
train_ds, val_ds = random_split(tmp_ds, [10000, 4034],  
                                generator=torch.Generator().manual_seed(42))  
test_ds = datasets.ImageFolder(test_dir, transform=test_transform)  
  
train_loader = torch.utils.data.DataLoader(train_ds, batch_size=64, shuffle=True)  
val_loader = torch.utils.data.DataLoader(val_ds, batch_size=64)  
test_loader = torch.utils.data.DataLoader(test_ds, batch_size=64)
```

```
In [8]: def show_batch(loader):  
        plt.figure(figsize=(60,60))  
        batch = next(iter(loader))  
        images, labels = batch  
        grid = make_grid(images, nrow = 10)  
        plt.imshow(np.transpose(grid, (1,2,0)))  
        plt.show()  
  
        show_batch(train_loader)
```



Choose Device

```
In [9]: device = torch.device('cuda' if torch.cuda.is_available else 'cpu')  
        device
```

```
Out[9]: device(type='cuda')
```

We will be using the pretrained resnet50 model.

```
In [10]: resnet = models.resnet50(pretrained=True)
# Freeze model params
for param in resnet.parameters():
    param.requires_grad = False
# Pull final fc layer feature dimensions
in_features = resnet.fc.in_features
out_features = resnet.fc.out_features
print(f"model in features: {in_features}")
print(f"model out features: {out_features}")

#We have to transform the 1000 outputs to 6 for our classification
classifier = nn.Sequential(OrderedDict([('fc1', nn.Linear(in_features, 512)),
                                       ('relu', nn.ReLU()),
                                       ('drop', nn.Dropout(0.05)),
                                       ('fc2', nn.Linear(512, 6)),
                                       ]))

# Now we add our classifier layer to our resnet model and then push it to our device
resnet.classifier = classifier
resnet.to(device)
```

Downloading: "https://download.pytorch.org/models/resnet50-0676ba61.pth" to /root/.cache/torch/hub/checkpoints/resnet50-0676ba61.pth

0%| | 0.00/97.8M [00:00<?, ?B/s]

model in features: 2048
model out features: 1000

```
Out[10]: ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  (layer1): Sequential(
    (0): Bottleneck(
      (conv1): Conv2d(64, 64, kernel_size=(1, 1), stride=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (conv3): Conv2d(64, 256, kernel_size=(1, 1), stride=(1, 1), bias=False)
```

Define our loss function, optimizer and scheduler (we use the multistep learning rate)

```
In [11]: criterion = nn.CrossEntropyLoss()
#initialize our learning rate to 0.001
optimizer = torch.optim.SGD(resnet.parameters(), lr=0.001)
scheduler = torch.optim.lr_scheduler.MultiStepLR(optimizer, milestones=[10, 15], gamma=0.05)
```

```
In [12]: epochs = 20

tr_losses = []
avg_epoch_tr_loss = []
tr_accuracy = []

val_losses = []
avg_epoch_val_loss = []
val_accuracy = []
val_loss_min = np.Inf

resnet.train()
for epoch in range(epochs):
    for i, batch in enumerate(train_loader):
        data, label = batch
        data, label = data.to(device), label.to(device)
        optimizer.zero_grad()
        logit = resnet(data)
        loss = criterion(logit, label)
        loss.backward()
        optimizer.step()
        tr_losses.append(loss.item())
        tr_accuracy.append(label.eq(logit.argmax(dim=1)).float().mean())
    print(f'\nEpoch No: {epoch + 1}, Training Loss: {torch.tensor(tr_losses).mean():.2f}, Training Accuracy: {torch.tensor(tr_ac
avg_epoch_tr_loss.append(torch.tensor(tr_losses).mean())

    resnet.eval()
    for i, batch in enumerate(val_loader):
        data, label = batch
        data, label = data.to(device), label.to(device)
        with torch.no_grad():
            logit = resnet(data)
            loss = criterion(logit, label)
            val_losses.append(loss.item())
            val_accuracy.append(label.eq(logit.argmax(dim=1)).float().mean())

    print(f'\nEpoch No: {epoch + 1}, Val Loss: {torch.tensor(val_losses).mean():.2f}, Val Accuracy: {torch.tensor(val_accuracy
avg_epoch_val_loss.append(torch.tensor(val_losses).mean())


    if torch.tensor(val_losses).float().mean() <= val_loss_min:
        torch.save(resnet.state_dict(), './model_state.pt')
        val_loss_min = torch.tensor(val_losses).mean()
    scheduler.step()
```

Epoch No: 1, Training Loss: 2.91, Training Accuracy: 0.50
Epoch No: 1, Val Loss: 0.87, Val Accuracy: 0.77
Epoch No: 2, Training Loss: 1.74, Training Accuracy: 0.65
Epoch No: 2, Val Loss: 0.73, Val Accuracy: 0.79
Epoch No: 3, Training Loss: 1.31, Training Accuracy: 0.72
Epoch No: 3, Val Loss: 0.66, Val Accuracy: 0.80
Epoch No: 4, Training Loss: 1.09, Training Accuracy: 0.75
Epoch No: 4, Val Loss: 0.61, Val Accuracy: 0.81
Epoch No: 5, Training Loss: 0.95, Training Accuracy: 0.77
Epoch No: 5, Val Loss: 0.57, Val Accuracy: 0.81
Epoch No: 6, Training Loss: 0.85, Training Accuracy: 0.79
Epoch No: 6, Val Loss: 0.54, Val Accuracy: 0.82
Epoch No: 7, Training Loss: 0.78, Training Accuracy: 0.80
Epoch No: 7, Val Loss: 0.54, Val Accuracy: 0.82
Epoch No: 8, Training Loss: 0.72, Training Accuracy: 0.81
Epoch No: 8, Val Loss: 0.53, Val Accuracy: 0.82
Epoch No: 9, Training Loss: 0.68, Training Accuracy: 0.82
Epoch No: 9, Val Loss: 0.51, Val Accuracy: 0.83
Epoch No: 10, Training Loss: 0.64, Training Accuracy: 0.82
Epoch No: 10, Val Loss: 0.50, Val Accuracy: 0.83
Epoch No: 11, Training Loss: 0.61, Training Accuracy: 0.83
Epoch No: 11, Val Loss: 0.49, Val Accuracy: 0.83
Epoch No: 12, Training Loss: 0.59, Training Accuracy: 0.83
Epoch No: 12, Val Loss: 0.48, Val Accuracy: 0.84
Epoch No: 13, Training Loss: 0.56, Training Accuracy: 0.84
Epoch No: 13, Val Loss: 0.47, Val Accuracy: 0.84
Epoch No: 14, Training Loss: 0.54, Training Accuracy: 0.84
Epoch No: 14, Val Loss: 0.46, Val Accuracy: 0.84
Epoch No: 15, Training Loss: 0.53, Training Accuracy: 0.85
Epoch No: 15, Val Loss: 0.45, Val Accuracy: 0.84
Epoch No: 16, Training Loss: 0.51, Training Accuracy: 0.85
Epoch No: 16, Val Loss: 0.45, Val Accuracy: 0.85
Epoch No: 17, Training Loss: 0.50, Training Accuracy: 0.85
Epoch No: 17, Val Loss: 0.44, Val Accuracy: 0.85
Epoch No: 18, Training Loss: 0.49, Training Accuracy: 0.85
Epoch No: 18, Val Loss: 0.44, Val Accuracy: 0.85
Epoch No: 19, Training Loss: 0.48, Training Accuracy: 0.86
Epoch No: 19, Val Loss: 0.43, Val Accuracy: 0.85
Epoch No: 20, Training Loss: 0.47, Training Accuracy: 0.86
Epoch No: 20, Val Loss: 0.43, Val Accuracy: 0.85

We see that we have reached a decent accuracy, and the val accuracy and training accuracy are really close, which means that the model isn't overfitting.

Plot of the average loss in each epoch

```
In [18]: plt.plot(avg_epoch_tr_loss, label='Training loss')
plt.plot(avg_epoch_val_loss, label='Validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```

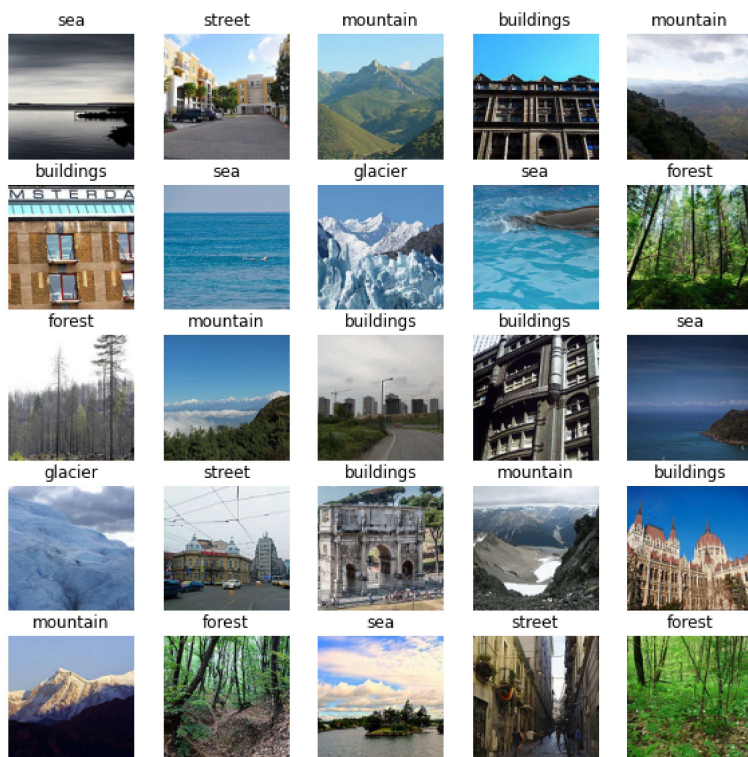


Now we make some predictions to test our model.

```
In [33]: resnet.load_state_dict(torch.load('./model_state.pt'))
resnet.eval()
#push our model to cuda
resnet = resnet.cuda()

def predictions(image):
    transform = test_transform(image)
    img = transform.unsqueeze(0).cuda()
    # push to gpu
    gpu_img = img.to(device)
    output = resnet(gpu_img)
    index = output.data.cpu().numpy().argmax()
    return index
```

```
In [36]: resnet.eval()
plt.figure(figsize=(10,10))
for i, images in enumerate(pred_files):
    if i > 24:
        break
    img = Image.open(images)
    index = predictions(img)
    plt.subplot(5,5,i+1)
    plt.title(classes[index])
    plt.axis('off')
    plt.imshow(img)
```



We see that all the predictions are correct and the model has a decent accuracy.

In []: